

Tanzim Hossain Romel

tanzimho@ualberta.ca | tanhrommel@gmail.com | tanzimhromel.com

LinkedIn: [thromel](#) | GitHub: [thromel](#) | Google Scholar

Edmonton, Alberta, Canada

University of Alberta M.Sc. student with about three years of professional software-engineering experience, building dependable AI-agent, developer-tooling, and healthcare analytics systems.

RESEARCH INTERESTS

AI4SE, AI for SRE, LLM4Coding, Trustworthy AI, Long-horizon Coding Agents

EDUCATION

• University of Alberta

M.Sc. in Computing Science

Sep 2026 - Present

Edmonton, Canada

- Thesis-based M.Sc. student in the [U-A-Goose](#) research group, supervised by [Dr. Zhou Yang](#), with affiliation to [Amii](#) (Alberta Machine Intelligence Institute).
- Current coursework: Human-AI Interaction and Formal Verification.

• Bangladesh University of Engineering and Technology (BUET)

B.Sc. in Computer Science and Engineering

Mar 2018 - May 2023

Dhaka, Bangladesh

- Selected coursework: Machine Learning, Operating Systems, Computer Security, Fault-Tolerant Systems, and High-Performance Database Systems.

RESEARCH EXPERIENCE

• University of Illinois Urbana-Champaign (UIUC)

Remote Research Intern through UIUC++ SRSE 2026; Professor Tianyin Xu's lab

Jun 2026 - Aug 2026

- Curated real-world postmortem reports into reproducible SRE benchmark problems on [SREGym](#).
- Worked on software and systems reliability research with emphasis on rigorous empirical evaluation and reproducible research artifacts.
- Prepared research prototypes, benchmark artifacts, and evaluation evidence for agentic software engineering and reliability tasks.
- Built SRE-style benchmark artifacts around fault injection, mitigation oracles, shortcut rejection, and lifecycle evidence for Kubernetes reliability scenarios.

• An Empirical Study on Remote Code Execution in Machine Learning Model Hosting Ecosystems

Tools: Python, Bandit, CodeQL, Semgrep, YARA | Under review at ICSE 2027; preprint: [arXiv:2601.14163](#)

Jun 2025 - Oct 2025

- Conducted a large-scale study of ~45,000 ML model repositories across Hugging Face, ModelScope, OpenCSG, OpenMMLab, and PyTorch Hub.
- Detected unsafe deserialization, eval-injection, malware-signature, and platform-safety risks using Bandit, CodeQL, Semgrep, YARA, and CWE-based classification.
- Analyzed 600+ developer discussions to build a taxonomy of misconceptions around remote code execution, trust flags, and model-loading safety.
- Compared platform defenses such as warning systems, scanning, SafeTensors support, and `trust_remote_code` controls to connect repository-level findings with ecosystem-level mitigation gaps.

• The Choice Can Be the Attack: Auditing Aligned Backdoors in LLM Agents

Tools: Python, vLLM, PyTorch, pandas, NumPy | Submitted to TACL 2026

Aug 2025 - Present

- Studied a hidden backdoor risk where a trigger changes which valid option an LLM agent chooses while the final answer still appears correct.
- Built **SHIFT**, which reruns the same task with and without a known trigger, then checks whether the changed choice favors the attacker's target rather than normal reasons like price or rating.
- Instrumented structured choice tasks with valid-option sets, logged option features, and matched no-trigger/trigger endpoint runs so the audit does not require model weights, logits, or hidden states.

• Multi-Agent Framework for Generating Relational DB Schema & ERD

Tools: Python, LangGraph, StateGraph, Z3 Solver, SQLAlchemy, Text2Schema

Jul 2025 - Present

- Extended SchemaAgent with Dr. Sukarna Barua (BUET) to turn natural-language requirements into relational schemas and ER diagrams.
- Built a LangGraph pipeline for entity extraction, relationship detection, normalization checks, Z3 validation, and targeted retry.
- Designed verification gates for entity/relation completeness, functional dependencies, normal-form compliance, key soundness, requirement coverage, and executable DDL repair.

SELECTED RESEARCH SOFTWARE AND SYSTEMS PROJECTS

• **ctxhelm + HelmBench: Context Compiler and Source-Free Agent Evaluation**

May 2026 - Present

Tools: Rust, MCP, CLI, Git, Hybrid Retrieval, Symbol Indexing, JSON Schema, Agent Tracing

- Built a local-first, read-only context compiler for coding agents, combining lexical search, symbol search, dependency expansion, related-test mapping, git co-change hints, semantic metadata, and memory signals.
- Shipped ctxhelm releases with archive/Homebrew installation, checksum verification, doctor checks, MCP smoke tests, and source-free evaluation artifacts.
- Built HelmBench as active evaluation infrastructure with trace schemas, privacy checks, evidence bundles, matrix runs, and quality gates that separate context-selection quality from downstream patch quality without exposing source, prompts, transcripts, or terminal logs.

• **Yet Another C Compiler**

Jun 2021 - Aug 2021; revised Oct 2024

Tools: C++17, Flex, Bison, CMake, x86-64 Assembly, SSA/IR Optimization

- Built a C compiler with lexer/parser, semantic analysis, x86-64 code generation, and a linear-scan register allocator.
- Modernized the Flex/Bison prototype with modular IR passes and SSA optimizations including SCCP, GVN, LICM, and dead-code elimination.

• **Bangla Handwritten Digit Recognition**

Jan 2023 - Feb 2023

Tools: Python, NumPy, OpenCV, CNN from Scratch, NumtaDB

- Implemented a NumPy CNN from scratch to recognize handwritten Bangla digits from the NumtaDB dataset.
- Built the preprocessing and training pipeline, reaching **95.9% accuracy** and placing **2nd among 120 students**.

OPEN SOURCE CONTRIBUTIONS

• **Selected Upstream Contributions**

2025 - 2026

Selected upstream contributions across SRE benchmarks, program analysis, .NET, web servers, agent systems, and compilers

[GitHub](#)

- **SREGym**: authored and hardened Kubernetes SRE benchmark work, including scheduler priority-preemption and Calico route-reflector label-drift scenarios with fault/oracle design, validation checks, and reviewer-facing evidence.
- **RefactoringMiner**: merged local MCP tooling, WebDiff workflow support, merge-parent-aware commit diffs, assertion migration matching, and single-page PR viewed toggles for refactoring-analysis workflows.
- **EF Core**: merged targeted data-stack changes for runtime migration creation/application, ON DELETE SET DEFAULT support, and nullable complex-property reload handling.
- **GenHTTP**: added automatic request-body decompression and binary-response support with focused framework behavior tests.
- **deepagents**: merged a CLI fix that stops and clears LoadingWidget animation timers on stop and unmount, preventing interval leaks in long-running terminal sessions and adding regression coverage for both cleanup paths.
- **TypeScript**: merged an upstream-aligned declaration-emit alias-resolution fix that preserves imported type aliases for inferred exports across module boundaries in Microsoft's Go port of TypeScript.

PROFESSIONAL EXPERIENCE

• **University of Alberta**

Sep 2026 - Present

Graduate Teaching and Research Assistant (GTA, GRA, GRAF)

Edmonton, Canada

- Teaching assistant for CMPUT 174.

- M.Sc. student supporting graduate teaching and research in software engineering and dependable AI with the [U-A-Goose](#) group.

• **IQVIA**

Jun 2023 - Jun 2026

Software Development Engineer

Dhaka, Bangladesh

- Built and maintained [KPI Library](#), the dynamic reporting layer within [IQVIA Orchestrated Analytics](#), using C#/.NET, EF Core, and AWS for healthcare analytics workflows.
- Refactored a complex healthcare-analytics filter path into reusable query components, reducing duplicated business rules and improving filtering performance by 70% in the targeted workflow.
- Reworked selected high-latency analytics paths with PostgreSQL query planning, targeted indexing, materialized views, Redis caching, and MongoDB tuning, reducing complex query time by 60% and API latency by 40%.
- Designed Playwright-backed export regression tests with multi-layer hashing to compare generated artifacts deterministically, raising export-workflow coverage from 72% to 95% and catching content-level regressions.
- Added OpenTelemetry and Jaeger tracing around API, query, and cache paths, making latency and regression sources visible across distributed services.
- Built LangGraph/LangChain workflows with MCP, RAG, and deepagents to map analyst prompts into structured dashboard setup, modification, drill-down analysis, and validated export actions, reducing setup effort by 85% in the targeted workflow.
- Optimized CI/CD build stages and parallel test execution, reducing deployment time by 35%; led POC work, mentored junior developers, and received the IQVIA Impact Program Silver Award for essential feature delivery.

SKILLS

Primary: Python, C#/.NET, Rust, TypeScript, Go, PostgreSQL, AWS, Docker, Kubernetes, LangGraph, MCP
Research and agent tooling: CodeQL, Semgrep, Bandit, YARA, Z3, RefactoringMiner, deepagents, vLLM, JSON Schema, source-free tracing
Systems and ML tooling: C++, Java, Flex, Bison, CMake, NS-3, NumPy, OpenCV, PyTorch, Linux namespaces
Backend and production tooling: EF Core, SQLAlchemy, REST APIs, Redis, MongoDB, SQL Server, DynamoDB, OpenTelemetry, Jaeger, Playwright

HONORS AND AWARDS

- **2nd Place – Bangla Handwritten Digits Recognition:** 95.9% accuracy, BUET ML Lab contest 2022
- **Dean’s List Award:** academic performance, BUET Level 2
- **National Science Olympiads:** prize winner in Bangladesh Physics and Chemistry Olympiads 2017
- **Talentpool HSC Scholarship:** 15th in Rajshahi Board with 95.6% marks 2017

REFERENCES

References available upon request.